

SANZHAR SOVET

Astana, Kazakhstan | +7 705 129 40 08 | sansan503451@gmail.com | [Website](#) | [LinkedIn](#) | [GitHub](#)

EDUCATION

Astana IT University <i>Master's degree in Applied Data Analytics — Full Grant, GPA 3.6</i>	Astana, Kazakhstan 2023 – 2025
Astana IT University <i>Bachelor's degree in Big Data Analysis — Full Grant, GPA 3.3</i>	Astana, Kazakhstan 2020 – 2023

EXPERIENCE

AI/ML Engineer LLMOps <i>Global Optima Technologies (founded by an ex-Meta/Amazon senior ML engineer)</i> - Maintained MLOps infrastructure across multiple GPU servers, covering deployment and monitoring of ASR, LLM, and TTS models (FastAPI, LiteLLM, Docker Swarm, Traefik, and Grafana, Prometheus) - Optimized runtime inference for latency-throughput trade-offs (vLLM, faster-whisper, Orpheus TTS, and Locust) - Fine-tuned LLMs for Kazakh language using LoRA/QLoRA methods (Unsloth) - Built AI agents for business process automation, with evaluation, tracing, and monitoring (LangChain/LangGraph and Langfuse)	Oct 2025 – Present
Junior Researcher <i>AI Research Center at Astana IT University</i> - Developed ML models for EEG-based attention monitoring as part of NeuroPro.kz - Authored publications on EEG-based cognitive states and ML methods in Q1/Q2 international journals - Collaborated on patent-pending AI developments, working with neuroscientists	Mar 2025 – Present

PROJECTS

Til-Qazyna ITutor <i>Docker Swarm, FastAPI, vLLM, Unsloth</i> - Configured and deployed LLM, ASR, and TTS APIs on 3 servers with an A10 GPU each, and set up monitoring and load balancing systems to serve 100 users - Optimized ASR, LLM, and TTS model inference, cutting median latency by 2x for ASR, 3x for LLM, and 4x for TTS with 10 concurrent users - Fine-tuned an LLM for the Kazakh language using LoRA/QLoRA, improving model quality by 10% on internal benchmarks	ITutor
Samruk-Energy AI Platform <i>LangGraph, LangChain, Langfuse, Docling</i> - Built an AI platform for business process automation, from LLM infrastructure and specialized agents with evaluation, tracing, and monitoring to the frontend, authentication, and integrations with external data sources - Developed a multi-agent system of three specialized agents * Generative BI (text-to-SQL) agent — TODO what it does * Accounting report generation agent — TODO what it does * Tender documentation analysis agent — TODO what it does	
Master's Thesis <i>NumPy, Pandas, Matplotlib, PyTorch, PyGame</i> - Published and won III place (certificate) at conference «2025 IEEE AITU: Digital Generation» - Built an evolutionary reinforcement learning model for decentralized multi-agent systems, where agents controlled by a Neural Network evolved through genetic algorithms to optimize survival in a 2D grid environment	GitHub

SKILLS

Programming: Python, SQL, Bash, HTML/CSS/JavaScript

Machine Learning & LLM: NumPy, Pandas, Scikit-learn, PyTorch, TensorFlow, Hugging Face Transformers, Unsloth

Agentic AI & RAG: LangGraph, LangChain, Pydantic, MCP, OpenAI API, Anthropic API, Qdrant, Docling, ASR (faster-whisper), Orpheus TTS

Infra & Serving: FastAPI, vLLM, LiteLLM, Docker, Docker Swarm, Kubernetes (kubect1), Traefik

MLOps & Observability: MLflow, W&B, Prometheus, Grafana, Locust, Langfuse, LangSmith, LLM-as-judge

Developer Tools: Git, CI/CD (GitLab CI), Claude Code, Codex, context engineering, custom skills/rules

Languages: English (B2), Russian (Native)